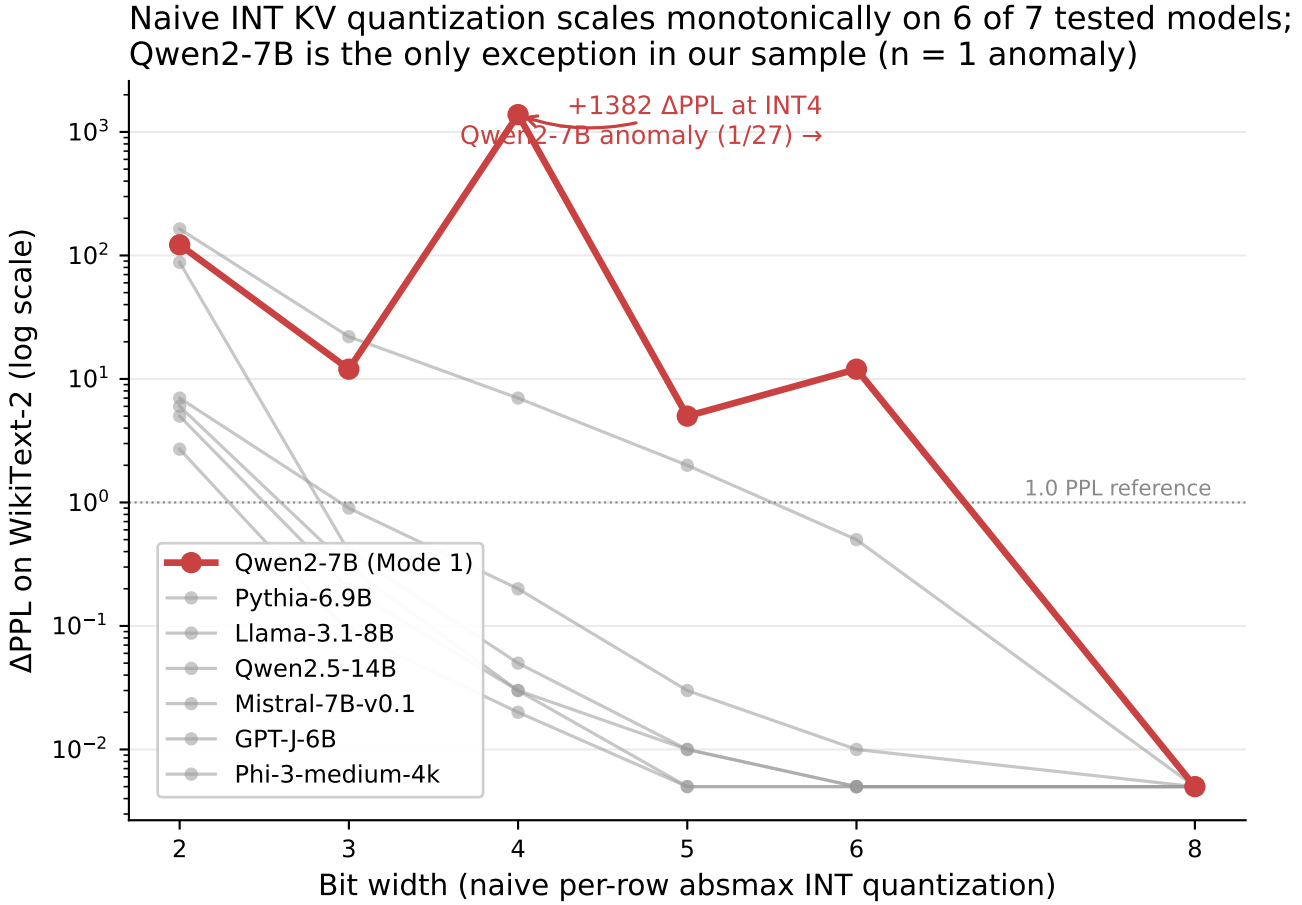
Naive INT KV quantization scales monotonically on 6 of 7 tested models;
Qwen2-7B is the only exception in our sample (n = 1 anomaly)
+1382 ΔPPL at INT4
Qwen2-7B anomaly (1/27) →
1.0 PPL reference
Qwen2-7B (Mode 1)
Pythia-6.9B
Llama-3.1-8B
Qwen2.5-14B
Mistral-7B-v0.1
GPT-J-6B
Phi-3-medium-4k
ΔPPL on WikiText-2 (log scale)
Bit width (naive per-row absmax INT quantization)